

Professional Summary

- **PhD Candidate in Computer Science**
at Yeoh's Optimization and Decision Analytics (YODA) Lab, Washington University in St. Louis, USA
- **Research Interests:** *Human-AI Interaction/ LLM Reasoning / Bandits & Reinforcement Learning*

Education

Washington University in St. Louis <i>PhD in Computer Science GPA: 3.84</i>	Saint Louis, Missouri, USA <i>Sep. 2023 – May 2027 (Expected)</i>
ShanghaiTech University (Co-Founded by the Chinese Academy of Sciences) <i>Master's in Information & Communication Engineering Major GPA: 3.7 Overall GPA: 3.66</i>	Shanghai, China <i>Sep. 2020 – Jun. 2023</i>
ShanghaiTech University <i>Bachelor's in Computer Science & Technology Major GPA: 3.5 Overall GPA: 3.43</i>	Shanghai, China <i>Sep. 2016 – Jun. 2020</i>

Research Experience

Human-AI Interaction Research Leader	Sep. 2023 – Present
<ul style="list-style-type: none"> • Dynamic and Personalized Probabilistic Human Modeling <ul style="list-style-type: none"> – Proposed Persona, a dynamic user modeling framework that shifts from static profiles to real-time adaptation via <i>Argumentation-based Dialogues</i>. – Integrated <i>Prospect Theory</i> (specifically probability weighting) with <i>Bayesian Inference</i> to mathematically capture evolving user beliefs and cognitive biases during interactions. – Demonstrated through human-subject evaluations that the framework significantly outperforms state-of-the-art methods in personalization and intent alignment. – Publication: Yinxu Tang, Stylianos Loukas Vasileiou & William Yeoh. “Does Your AI Agent Get You? A Personalizable Framework for Approximating Human Models from Argumentation-based Dialogue Traces.” AAAI Conference on Artificial Intelligence, 2025 (Oral). 	
<ul style="list-style-type: none"> • Probabilistic Model Reconciliation Framework <ul style="list-style-type: none"> – Pioneered the first model reconciliation framework grounded in <i>Probabilistic Logic Programming (PLP)</i> to address uncertainty in human knowledge, moving beyond traditional deterministic approaches[cite: 30, 36]. – Formulated explanations as <i>Cost-Optimal Model Updates</i> that reconcile inconsistencies in <i>Maximum Posteriori Estimation (MPE)</i> outcomes between AI and human models. – Engineered optimized search algorithms (including <i>Greedy</i> and <i>Weighted</i> variants) guided by theoretical structural insights, achieving significant scalability and efficiency over generic baselines. – Validated the framework’s effectiveness through human-subject studies on explanation preferences and computational experiments demonstrating superior performance. – Publication: Yinxu Tang, Stylianos Loukas Vasileiou & William Yeoh. “Model Reconciliation via Cost-Optimal Explanations in Probabilistic Logic Programming” Neural Information Processing Systems (NeurIPS), 2025 (poster). 	
LLM Reasoning Research Leader	Aug. 2025 – Present
<ul style="list-style-type: none"> • LLM Reasoning for Knowledge Graph Question Answering (UniRel-R1) <ul style="list-style-type: none"> – Formalized the novel task of <i>Relation-Centric KGQA</i>, shifting the paradigm from single-entity prediction to generating explanatory subgraphs that capture semantic associations. – Developed UniRel-R1, a unified framework integrating multi-stage graph pruning with an <i>RL-tuned LLM</i> to navigate vast search spaces and identify unique reasoning paths. – Engineered a custom reward function that incentivizes <i>compactness, specificity, and informativeness</i>. – Achieved significant gains in connectivity and answer quality over baselines, demonstrating strong generalization capabilities to unseen entities and relations. – Relavent Paper: Yinxu Tang, Chengsong Huang, Jiaxin Huang & William Yeoh “UniRel-R1: RL-tuned LLM Reasoning for Knowledge Graph Relational Question Answering.” (Under Review) 	
Bandits & Reinforcement Learning	Feb. 2019 – Jun. 2023
<ul style="list-style-type: none"> • Constrained Bandits Research Leader and Active Contributor <ul style="list-style-type: none"> – Formulated a <i>Stochastic Optimization</i> framework for Multi-Armed Bandits (MAB) with non-linear objective functions and long-term constraints. 	

- Derived theoretical Sub-linear Regret Bounds and constraint violation guarantees using Lyapunov Optimization and Upper Confidence Bound (UCB) techniques.
- Implemented and validated predictive models using BERT, LSTM, and RNNs for high-dimensional sequence prediction, analyzing convergence rates in edge intelligence systems.
- **Publications:**
 - (a) **Yinxu Tang**, Jianfeng Hou, Xi Huang, Ziyu Shao & Yang Yang. “*Green Edge Intelligence Scheme for Mobile Keyboard Emoji Prediction.*” IEEE Transactions on Mobile Computing (TMC), 2024.
 - (b) Xi Huang, **Yinxu Tang**, Ziyu Shao & Yang Yang. “*Joint Switch-Controller Association & Control Devolution for SDN Systems: An Integrated Online Perspective of Control & Learning.*” IEEE Transactions on Network & Management (TNSM), 2021.
 - (c) Jianfeng Hou, **Yinxu Tang**, Xi Huang, Ziyu Shao & Yang Yang. “*Green Edge Intelligence Scheme for Mobile Keyboard Emoji Prediction.*” IEEE International Conference on Communications (ICC), 2021.
 - (d) Xin Gao, Xi Huang, **Yinxu Tang**, Ziyu Shao & Yang Yang. “*History-Aware Online Cache Placement in Fog-Assisted IoT Systems: An Integration of Learning & Control.*” IEEE Internet of Things Journal (IoT-J), 2021.
 - (e) Xi Huang, **Yinxu Tang**, Ziyu Shao & Yang Yang. “*Joint Switch-Controller Association & Control Devolution for SDN Systems: An Integration of Online Control & Online Learning.*” IEEE/ACM International Symposium on Quality of Service (IWQoS), 2020.
 - (f) Xin Gao, Xi Huang, **Yinxu Tang**, Ziyu Shao & Yang Yang. “*Proactive Cache Placement with Bandit Learning in Fog-Assisted IoT Systems.*” IEEE International Conference on Communications (ICC), 2020.
 - (g) Junge Zhu, Xi Huang, **Yinxu Tang** & Ziyu Shao. “*Learning-Aided Online Task Offloading for UAVs-Aided IoT Systems.*” IEEE Vehicular Technology Conference (VTC), 2019.
- **Constrained Graphical Bandits | Active Contributor**
 - Applied online learning algorithms that exploit Graph Feedback Mechanisms (side-observations) to accelerate convergence and reduce exploration variance.
 - Conducted rigorous theoretical analysis to characterize the dependency of Regret Bounds on graph topology metrics, specifically the Independence Number and Degree Centrality.
 - Performed extensive simulations on diverse graph structures to benchmark algorithm robustness and quantify the trade-off between information flow and learning efficiency.
- **Publications:**
 - (a) Shangshang Wang, Simeng Bian, **Yinxu Tang** & Ziyu Shao. “*Social-Aware Distributed Meta-Learning: A Perspective of Constrained Graphical Bandits.*” IEEE International Conference on Communications (ICC), 2023.
 - (b) Simeng Bian, Shangshang Wang, **Yinxu Tang** & Ziyu Shao. “*Social-Aware Edge Intelligence: A Constrained Graphical Bandit Approach.*” IEEE Global Communications Conference (GLOBECOM), 2022.
- **Privacy-Preserving Constrained Bandits | Research Leader**
 - Incorporated Local Differential Privacy (ϵ -LDP) mechanisms to secure distributed learning, establishing rigorous privacy guarantees via noise injection.
 - Derived theoretical bounds to quantify the Privacy-Utility Trade-off, mathematically characterizing how the privacy budget (ϵ) impacts the algorithm’s Regret Convergence.
 - Validated the theoretical analysis through extensive simulations, demonstrating algorithm robustness under varying noise regimes and privacy constraints.
 - **Publication:** Tianyi Zhang, Shangshang Wang, **Yinxu Tang**, Ziyu Shao & Yang Yang. “*Privacy-Preserving Edge Intelligence: A Perspective of Constrained Bandits.*” IEEE Wireless Communications and Networking Conference (WCNC), 2024.
- **Bandits with Nash Equilibrium | Research Leader**
 - Formulated a decentralized resource allocation framework integrating Multi-Armed Bandits with the Gale-Shapley Deferred Acceptance mechanism to solve two-sided matching problems.
 - Proved theoretical convergence to Nash Stability (Stable Matching) and analyzed the trade-off between exploration regret and matching stability in dynamic environments.
 - **Publication:** **Yinxu Tang**, Tao Huang, Xi Huang, Ziyu Shao & Yang Yang. “*Learning-Aided Stable Matching for Switch Controller Association in SDN Systems.*” IEEE International Conference on Communications (ICC), 2022.

Networks | Active Contributor

Nov. 2023 – Mar. 2024

- Proposed a 6G-based deployment framework for RAG-enhanced generative services, emphasizing real-time knowledge base updates, service customization, and edge intelligence integration.
- Addressed key challenges using techniques such as data fusion, dynamic KB distribution, service customization, and user-interaction optimization in 6G environments.
- **Relevant Paper:**
 - (a) Xi Huang, **Yinxu Tang**, Junling Li, Ning Zhang & Xuemin Shen. “*Toward Effective Retrieval Augmented Generative Services in 6G Networks.*” IEEE Network, 2024.

Technical Skills

Languages: Python (PyTorch, TensorFlow, Scikit-learn, NumPy, Pandas), C++, C, Matlab.

Awards & Honors

- | | |
|---|-----------|
| • National Scholarship for Graduate Students — Ministry of Education of China | Oct. 2022 |
| • Merit Student — ShanghaiTech University | Dec. 2021 |
| • Meritorious Winner — Mathematical Contest in Modelling | Jun. 2018 |

Teaching

Probability & Mathematical Statistics

Fall 2022 & Fall 2021

Teaching Assistant

ShanghaiTech University

- Led weekly tutorial sessions for 70-80 students, focusing on exercises, discussions, and interactive learning.
- Designed and graded assignments, including weekly exercises, project proposals, and exam papers.
- Provided mentorship to students through one-on-one consultations and tailored guidance during office hours.